

Credit_Card_Fraud_Detection

January 30, 2025

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: df = pd.read_csv("creditcard.csv")
df.head()
```

```
[2]:
```

	Time	V1	V2	V3	V4	V5	V6	V7 \	V8	V9	...	V21	V22	V23	V24	V25 \
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

```
[3]: pd.options.display.max_columns = None
```

```
[4]: df.tail()
```

```
[4]:
```

	Time	V1	V2	V3	V4	V5 \
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961

```
284806 172792.0 -0.533413 -0.189733 0.703337 -0.506271 -0.012546
```

	V6	V7	V8	V9	V10	V11	V12	\
284802	-2.606837	-4.918215	7.305334	1.914428	4.356170	-1.593105	2.711941	
284803	1.058415	0.024330	0.294869	0.584800	-0.975926	-0.150189	0.915802	
284804	3.031260	-0.296827	0.708417	0.432454	-0.484782	0.411614	0.063119	
284805	0.623708	-0.686180	0.679145	0.392087	-0.399126	-1.933849	-0.962886	
284806	-0.649617	1.577006	-0.414650	0.486180	-0.915427	-1.040458	-0.031513	

	V13	V14	V15	V16	V17	V18	V19	\
284802	-0.689256	4.626942	-0.924459	1.107641	1.991691	0.510632	-0.682920	
284803	1.214756	-0.675143	1.164931	-0.711757	-0.025693	-1.221179	-1.545556	
284804	-0.183699	-0.510602	1.329284	0.140716	0.313502	0.395652	-0.577252	
284805	-1.042082	0.449624	1.962563	-0.608577	0.509928	1.113981	2.897849	
284806	-0.188093	-0.084316	0.041333	-0.302620	-0.660377	0.167430	-0.256117	

	V20	V21	V22	V23	V24	V25	V26	\
284802	1.475829	0.213454	0.111864	1.014480	-0.509348	1.436807	0.250034	
284803	0.059616	0.214205	0.924384	0.012463	-1.016226	-0.606624	-0.395255	
284804	0.001396	0.232045	0.578229	-0.037501	0.640134	0.265745	-0.087371	
284805	0.127434	0.265245	0.800049	-0.163298	0.123205	-0.569159	0.546668	
284806	0.382948	0.261057	0.643078	0.376777	0.008797	-0.473649	-0.818267	

	V27	V28	Amount	Class
284802	0.943651	0.823731	0.77	0
284803	0.068472	-0.053527	24.79	0
284804	0.004455	-0.026561	67.88	0
284805	0.108821	0.104533	10.00	0
284806	-0.002415	0.013649	217.00	0

```
[5]: df.shape
```

```
[5]: (284807, 31)
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype
---  -
0    Time    284807 non-null    float64
1    V1       284807 non-null    float64
2    V2       284807 non-null    float64
3    V3       284807 non-null    float64
4    V4       284807 non-null    float64
5    V5       284807 non-null    float64
6    V6       284807 non-null    float64
```

```

7   V7      284807 non-null float64
8   V8      284807 non-null float64
9   V9      284807 non-null float64
10  V10     284807 non-null float64
11  V11     284807 non-null float64
12  V12     284807 non-null float64
13  V13     284807 non-null float64
14  V14     284807 non-null float64
15  V15     284807 non-null float64
16  V16     284807 non-null float64
17  V17     284807 non-null float64
18  V18     284807 non-null float64
19  V19     284807 non-null float64
20  V20     284807 non-null float64
21  V21     284807 non-null float64
22  V22     284807 non-null float64
23  V23     284807 non-null float64
24  V24     284807 non-null float64
25  V25     284807 non-null float64
26  V26     284807 non-null float64
27  V27     284807 non-null float64
28  V28     284807 non-null float64
29  Amount  284807 non-null float64
30  Class   284807 non-null int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```

[7]: ## make the dataset to a standarized format
from sklearn.preprocessing import StandardScaler
import pandas as pd

# Assuming df is your DataFrame and it has a column 'Amount'
sc = StandardScaler()

# Correcting the use of DataFrame to standardize the 'Amount' column
df["Amount"] = sc.fit_transform(df[["Amount"]])

df.head()

```

```

[7]:   Time      V1      V2      V3      V4      V5      V6      V7  \
0   0.0 -1.359807 -0.072781  2.536347  1.378155 -0.338321  0.462388  0.239599
1   0.0  1.191857  0.266151  0.166480  0.448154  0.060018 -0.082361 -0.078803
2   1.0 -1.358354 -1.340163  1.773209  0.379780 -0.503198  1.800499  0.791461
3   1.0 -0.966272 -0.185226  1.792993 -0.863291 -0.010309  1.247203  0.237609
4   2.0 -1.158233  0.877737  1.548718  0.403034 -0.407193  0.095921  0.592941

      V8      V9      V10      V11      V12      V13      V14  \

```

0	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.991390	-0.311169
1	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.489095	-0.143772
2	0.247676	-1.514654	0.207643	0.624501	0.066084	0.717293	-0.165946
3	0.377436	-1.387024	-0.054952	-0.226487	0.178228	0.507757	-0.287924
4	-0.270533	0.817739	0.753074	-0.822843	0.538196	1.345852	-1.119670

	V15	V16	V17	V18	V19	V20	V21 \
0	1.468177	-0.470401	0.207971	0.025791	0.403993	0.251412	-0.018307
1	0.635558	0.463917	-0.114805	-0.183361	-0.145783	-0.069083	-0.225775
2	2.345865	-2.890083	1.109969	-0.121359	-2.261857	0.524980	0.247998
3	-0.631418	-1.059647	-0.684093	1.965775	-1.232622	-0.208038	-0.108300
4	0.175121	-0.451449	-0.237033	-0.038195	0.803487	0.408542	-0.009431

	V22	V23	V24	V25	V26	V27	V28 \
0	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053
1	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724
2	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752
3	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458
4	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153

	Amount	Class
0	0.244964	0
1	-0.342475	0
2	1.160686	0
3	0.140534	0
4	-0.073403	0

```
[8]: df.drop(["Time"],axis = 1)
```

```
[8]:
```

	V1	V2	V3	V4	V5	V6 \
0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388
1	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361
2	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499
3	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203
4	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921
...
284802	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837
284803	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415
284804	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260
284805	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708
284806	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617

	V7	V8	V9	V10	V11	V12	V13 \
0	0.239599	0.098698	0.363787	0.090794	-0.551600	-0.617801	-0.991390
1	-0.078803	0.085102	-0.255425	-0.166974	1.612727	1.065235	0.489095
2	0.791461	0.247676	-1.514654	0.207643	0.624501	0.066084	0.717293
3	0.237609	0.377436	-1.387024	-0.054952	-0.226487	0.178228	0.507757

4	0.592941	-0.270533	0.817739	0.753074	-0.822843	0.538196	1.345852
...
284802	-4.918215	7.305334	1.914428	4.356170	-1.593105	2.711941	-0.689256
284803	0.024330	0.294869	0.584800	-0.975926	-0.150189	0.915802	1.214756
284804	-0.296827	0.708417	0.432454	-0.484782	0.411614	0.063119	-0.183699
284805	-0.686180	0.679145	0.392087	-0.399126	-1.933849	-0.962886	-1.042082
284806	1.577006	-0.414650	0.486180	-0.915427	-1.040458	-0.031513	-0.188093

	V14	V15	V16	V17	V18	V19	V20 \
0	-0.311169	1.468177	-0.470401	0.207971	0.025791	0.403993	0.251412
1	-0.143772	0.635558	0.463917	-0.114805	-0.183361	-0.145783	-0.069083
2	-0.165946	2.345865	-2.890083	1.109969	-0.121359	-2.261857	0.524980
3	-0.287924	-0.631418	-1.059647	-0.684093	1.965775	-1.232622	-0.208038
4	-1.119670	0.175121	-0.451449	-0.237033	-0.038195	0.803487	0.408542
...
284802	4.626942	-0.924459	1.107641	1.991691	0.510632	-0.682920	1.475829
284803	-0.675143	1.164931	-0.711757	-0.025693	-1.221179	-1.545556	0.059616
284804	-0.510602	1.329284	0.140716	0.313502	0.395652	-0.577252	0.001396
284805	0.449624	1.962563	-0.608577	0.509928	1.113981	2.897849	0.127434
284806	-0.084316	0.041333	-0.302620	-0.660377	0.167430	-0.256117	0.382948

	V21	V22	V23	V24	V25	V26	V27 \
0	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558
1	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983
2	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353
3	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723
4	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422
...
284802	0.213454	0.111864	1.014480	-0.509348	1.436807	0.250034	0.943651
284803	0.214205	0.924384	0.012463	-1.016226	-0.606624	-0.395255	0.068472
284804	0.232045	0.578229	-0.037501	0.640134	0.265745	-0.087371	0.004455
284805	0.265245	0.800049	-0.163298	0.123205	-0.569159	0.546668	0.108821
284806	0.261057	0.643078	0.376777	0.008797	-0.473649	-0.818267	-0.002415

	V28	Amount	Class
0	-0.021053	0.244964	0
1	0.014724	-0.342475	0
2	-0.059752	1.160686	0
3	0.061458	0.140534	0
4	0.215153	-0.073403	0
...
284802	0.823731	-0.350151	0
284803	-0.053527	-0.254117	0
284804	-0.026561	-0.081839	0
284805	0.104533	-0.313249	0
284806	0.013649	0.514355	0

[284807 rows x 30 columns]

```
[9]: df.duplicated().any()
```

[9]: True

```
[10]: df = df.drop_duplicates()
```

```
[11]: df.shape
```

[11]: (283726, 31)

```
[12]: # Checking for imbalance
df["Class"].value_counts()
```

[12]: Class
0 283253
1 473
Name: count, dtype: int64

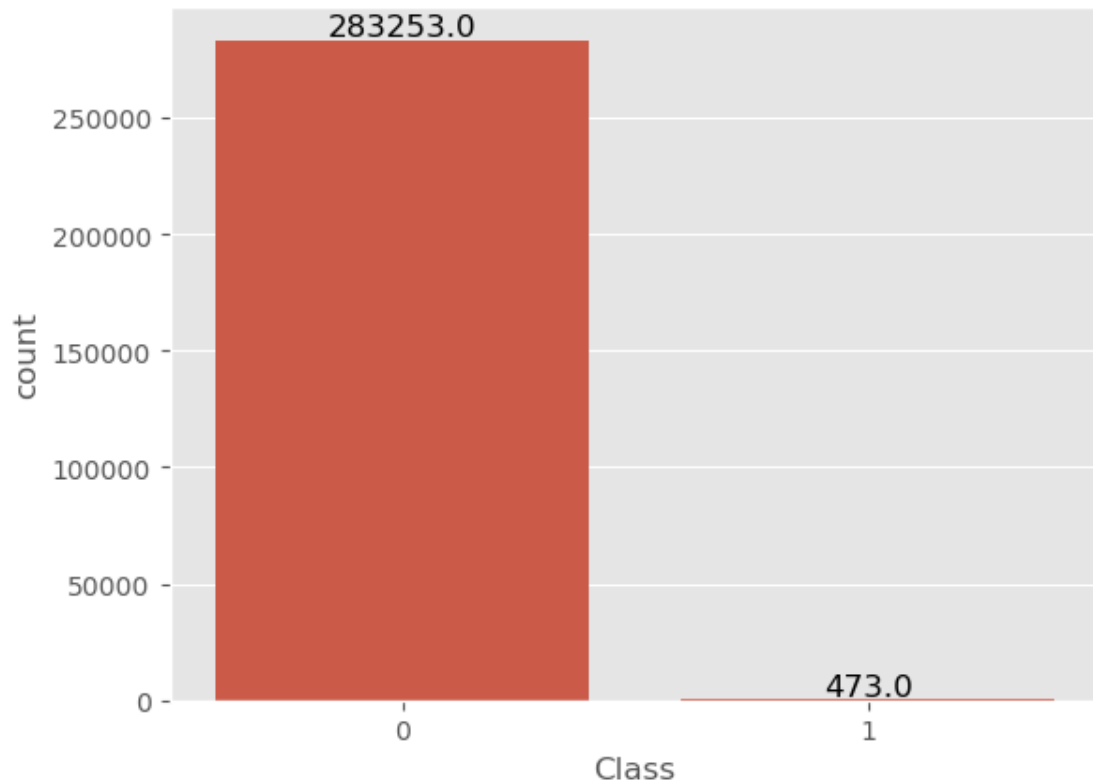
```
[13]: import seaborn as sns
import matplotlib.pyplot as plt
plt.style.use('ggplot')
```

```
[14]: import seaborn as sns
import matplotlib.pyplot as plt

# Create the countplot
ax = sns.countplot(x=df['Class'])

# Add data labels
for p in ax.patches:
    ax.annotate(f'{p.get_height()}',
                (p.get_x() + p.get_width() / 2., p.get_height()),
                ha='center', va='center',
                fontsize=12, color='black',
                xytext=(0, 5), textcoords='offset points')

plt.show()
```



```
[15]: X = df.drop('Class',axis=1)
      Y = df['Class']
```

```
[19]: from sklearn.model_selection import train_test_split
```

```
[20]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.
      ↪2,random_state = 42)
```

```
[21]: import numpy as np
      from sklearn.linear_model import LogisticRegression
      from sklearn.ensemble import RandomForestClassifier
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import accuracy_score, f1_score, precision_score,
      ↪recall_score
```

```
[22]: classifier = {
      "Logistic Regression" : LogisticRegression(),
      "Decision Tree Classifier" : DecisionTreeClassifier()
      }

      for name,clf in classifier.items():
```

```

print(f"\n===== {name} =====")
clf.fit(X_train,Y_train)
y_pred = clf.predict(X_test)
accuracy = accuracy_score(Y_test,y_pred)
print(f"\n Accuracy: {accuracy_score(Y_test,y_pred)}")
print(f"\n Precision: {precision_score(Y_test,y_pred)}")
print(f"\n Recall: {recall_score(Y_test,y_pred)}")
print(f"\n F1 Score: {f1_score(Y_test,y_pred)}")

```

=====Logistic Regression=====

C:\Users\afros\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Accuracy: 0.999048391076023

Precision: 0.7647058823529411

Recall: 0.5777777777777777

F1 Score: 0.6582278481012658

=====Decision Tree Classifier=====

Accuracy: 0.9990307686885419

Precision: 0.6842105263157895

Recall: 0.7222222222222222

F1 Score: 0.7027027027027027

```

[23]: ## Checking with Under sampling to match the lower side data
normal = df[df['Class'] == 0]
fraud = df[df['Class'] == 1]

```

```

[24]: normal_sample = normal.sample(n=473)

```



```
[25]: normal_sample.shape
```

```
[25]: (473, 31)
```

```
[26]: new_under_Sampled = pd.concat([normal_sample,fraud],ignore_index = True)
```

```
[27]: new_under_Sampled
```

```
[27]:
```

	Time	V1	V2	V3	V4	V5	V6	\
0	41639.0	1.161922	-0.476883	1.091284	0.211155	-1.094138	0.063974	
1	134971.0	-0.100642	1.234425	0.024438	1.032598	0.744236	-0.360728	
2	159612.0	2.035526	-0.085170	-1.173378	0.211055	0.144427	-0.606359	
3	29024.0	-0.407065	0.788201	1.551483	1.181414	0.202893	-0.205186	
4	2185.0	-0.553649	0.700380	1.134316	-0.070443	0.377868	1.762290	
..	
941	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	
942	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	
943	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	
944	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	
945	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	

	V7	V8	V9	V10	V11	V12	V13	\
0	-0.786758	0.242821	1.336610	-0.471299	-0.724728	0.186454	-0.937451	
1	0.608881	-0.127264	-0.520564	-0.055720	-0.458310	-0.390206	0.138354	
2	0.087886	-0.164929	0.238581	0.222041	0.786657	1.264790	0.457782	
3	0.449064	0.081953	-0.578149	-0.142392	-0.410134	-0.478763	-0.682878	
4	-0.094326	-0.232484	-0.766323	-0.147612	2.400641	0.905801	-0.151857	
..	
941	-0.882850	0.697211	-2.064945	-5.587794	2.115795	-5.417424	-1.235123	
942	-1.413170	0.248525	-1.127396	-3.232153	2.858466	-3.096915	-0.792532	
943	-2.234739	1.210158	-0.652250	-3.463891	1.794969	-2.775022	-0.418950	
944	-2.208002	1.058733	-1.632333	-5.245984	1.933520	-5.030465	-1.127455	
945	0.223050	-0.068384	0.577829	-0.888722	0.491140	0.728903	0.380428	

	V14	V15	V16	V17	V18	V19	V20	\
0	-0.348797	0.047390	-0.370529	0.474205	-0.978150	0.112303	-0.160556	
1	-0.796478	1.465520	-0.835594	1.467085	0.175582	2.762432	0.395922	
2	0.373220	-0.635767	0.108402	-0.626169	-0.408307	0.484958	-0.183367	
3	0.462423	1.613420	-0.548843	0.198135	-0.125325	0.293668	0.034179	
4	0.661012	2.116675	-1.575973	1.366089	-1.700175	0.667103	-0.109729	
..	
941	-6.665177	0.401701	-2.897825	-4.570529	-1.315147	0.391167	1.252967	
942	-5.210141	-0.613803	-2.155297	-3.267116	-0.688505	0.737657	0.226138	
943	-4.057162	-0.712616	-1.603015	-5.035326	-0.507000	0.266272	0.247968	
944	-6.416628	0.141237	-2.549498	-4.614717	-1.478138	-0.035480	0.306271	
945	-1.948883	-0.832498	0.519436	0.903562	1.197315	0.593509	-0.017652	

	V21	V22	V23	V24	V25	V26	V27 \
0	-0.214742	-0.372933	0.118152	0.136635	0.032242	0.974516	-0.023640
1	-0.299127	-0.620174	-0.008826	0.604631	-0.705073	0.583574	0.293141
2	-0.248886	-0.588268	0.289497	-0.383927	-0.285043	0.202634	-0.068878
3	0.189416	0.534401	-0.086630	0.056051	-0.295431	-0.226073	0.171159
4	0.820552	0.083399	0.270325	-0.975188	-1.111250	1.132361	0.208529
..
941	0.778584	-0.319189	0.639419	-0.294885	0.537503	0.788395	0.292680
942	0.370612	0.028234	-0.145640	-0.081049	0.521875	0.739467	0.389152
943	0.751826	0.834108	0.190944	0.032070	-0.739695	0.471111	0.385107
944	0.583276	-0.269209	-0.456108	-0.183659	-0.328168	0.606116	0.884876
945	-0.164350	-0.295135	-0.072173	-0.450261	0.313267	-0.289617	0.002988

	V28	Amount	Class
0	0.009608	-0.307331	0
1	0.256360	-0.346073	0
2	-0.073233	-0.349271	0
3	0.150457	-0.277666	0
4	0.115114	-0.198503	0
..
941	0.147968	1.206024	1
942	0.186637	-0.350191	1
943	0.194361	-0.041818	1
944	-0.253700	0.626302	1
945	-0.015309	-0.183191	1

[946 rows x 31 columns]

```
[29]: new_under_Sampled['Class'].value_counts()

## Problem for this is we have lost around 250000 data record
## so I would suggest to go for over sampling
```

```
[29]: Class
0      473
1      473
Name: count, dtype: int64
```

```
[30]: X = new_under_Sampled.drop('Class',axis=1)
Y = new_under_Sampled['Class']
```

```
[31]: X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size = 0.
↪2,random_state = 42)
```

```
[32]: classifier = {
        "Logistic Regression" : LogisticRegression(),
        "Decision Tree Classifier" : DecisionTreeClassifier()
```

```

}

for name,clf in classifier.items():
    print(f"\n===== {name} =====")
    clf.fit(X_train,Y_train)
    y_pred = clf.predict(X_test)
    accuracy = accuracy_score(Y_test,y_pred)
    print(f"\n Accuracy: {accuracy_score(Y_test,y_pred)}")
    print(f"\n Precision: {precision_score(Y_test,y_pred)}")
    print(f"\n Recall: {recall_score(Y_test,y_pred)}")
    print(f"\n F1 Score: {f1_score(Y_test,y_pred)}")

## Problem for this is we have lost around 250000 data record
## so I would suggest to go for over sampling

```

=====Logistic Regression=====

Accuracy: 0.9526315789473684

Precision: 0.979381443298969

Recall: 0.9313725490196079

F1 Score: 0.9547738693467337

=====Decision Tree Classifier=====

Accuracy: 0.9052631578947369

Precision: 0.92

Recall: 0.9019607843137255

F1 Score: 0.9108910891089109

C:\Users\afros\anaconda3\Lib\site-packages\sklearn\linear_model_logistic.py:469: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

        n_iter_i = _check_optimize_result(
[33]: ## Now lets try for over sampling lets use SMOT

X = df.drop('Class',axis=1)
Y = df['Class']

[34]: from imblearn.over_sampling import SMOTE

[35]: X_res ,Y_res = SMOTE().fit_resample(X,Y)

[37]: Y_res.value_counts()

[37]: Class
0      283253
1      283253
Name: count, dtype: int64

[38]: X_train,X_test,Y_train,Y_test = train_test_split(X_res,Y_res,test_size = 0.
↪2,random_state = 42)

[39]: classifier = {
        "Logistic Regression" : LogisticRegression(),
        "Decision Tree Classifier" : DecisionTreeClassifier()
    }

    for name,clf in classifier.items():
        print(f"\n===== {name} =====")
        clf.fit(X_train,Y_train)
        y_pred = clf.predict(X_test)
        accuracy = accuracy_score(Y_test,y_pred)
        print(f"\n Accuracy: {accuracy_score(Y_test,y_pred)}")
        print(f"\n Precision: {precision_score(Y_test,y_pred)}")
        print(f"\n Recall: {recall_score(Y_test,y_pred)}")
        print(f"\n F1 Score: {f1_score(Y_test,y_pred)}")

## Problem for this is we have lost around 250000 data record
## so I would suggest to go for over sampling

```

```

=====Logistic Regression=====

```

```

C:\Users\afros\anaconda3\Lib\site-
packages\sklearn\linear_model\_logistic.py:469: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

Accuracy: 0.9731160968032339

Precision: 0.9837761030271418

Recall: 0.9622794208202115

F1 Score: 0.9729090326941549

=====Decision Tree Classifier=====

Accuracy: 0.9982789359411132

Precision: 0.9974182444061962

Recall: 0.999155509421348

F1 Score: 0.9982861210965309

```
[40]: dtc = DecisionTreeClassifier()  
      dtc.fit(X_res,Y_res)
```

```
[40]: DecisionTreeClassifier()
```

```
[42]: import joblib  
  
      # Assuming dtc is your trained DecisionTreeClassifier model  
      joblib.dump(dtc, 'credit_card_model.pkl')
```

```
[42]: ['credit_card_model.pkl']
```

```
[43]: model = joblib.load("credit_card_model.pkl")
```

```
[45]: # Assuming `model` is your trained model and you're predicting on one row  
  
      pred = model.predict([[  
          -1.3598071336738, -0.0727811733098497, 2.53634673796914, 1.37815522427443,  
          -0.338320769942518, 0.462387777762292, 0.239598554061257, 0.  
          ↪0986979012610507,
```

```

    0.363786969611213, 0.0907941719789316, -0.551599533260813, -0.
↪617800855762348,
    -0.991389847235408, -0.311169353699879, 1.46817697209427, -0.
↪470400525259478,
    0.207971241929242, 0.0257905801985591, 0.403992960255733, 0.251412098239705,
    -0.018306777944153, 0.277837575558899, -0.110473910188767, 0.
↪0669280749146731,
    0.128539358273528, -0.189114843888824, 0.133558376740387, -0.
↪0210530534538215,
    149.62, # Make sure all features are included, like V29, V30, if any
    0.5      # Example of another missing feature, just as an illustration
])

```

C:\Users\afros\anaconda3\Lib\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but DecisionTreeClassifier was fitted with feature names

```
warnings.warn(
```

```
[46]: pred
```

```
[46]: array([0], dtype=int64)
```

```

[47]: if pred == 0:
        print("Normal Transaction")
    else:
        print("Fraud Transaction")

```

Normal Transaction